2025 IEEE International Symposium on Future Telecommunication Technologies (SOFTT) | 979-8-3315-6931-0/25/$31.00 ©2025 IEEE | DOI: 10.1109/SOFTT67007.2025.11213027

Designing an AIOps (Artificial Intelligent Operations) Framework for Energy-Efficient IT Operations: A Case Study on Cloud Computing Environment

Milankumar Rana
Information Technology
University of the Cumberlands
Williamsburg, USA
0009-0009-1623-0429

Sakhita Sree Gadde
Information Technology
Independent Researcher
College Station, USA
0009-0009-3675-1299

Abstract- **This study looks at using AI and quantum computing to increase the security and efficiency of financial transactions. The suggested system's hybrid architecture enables autonomous agents to work together to make decisions. Computer performance is enhanced and accelerated by quantum techniques. Grover's algorithm speeds up the process of finding the best solution and confirming transactions. Quantum key distribution (QKD) protocols are designed to be impenetrable and protect data. Agent-based utility models are even more flexible and scalable because they can adapt to the needs of the operational environment. The design places a high value on adaptability, ease of installation across multiple platforms, and compatibility with other systems. It has a cloud-friendly design. Real-world testing demonstrates that the new techniques are far more effective than the previous ones. They reduced transaction throughput by 64%, vulnerability exposure by 89%, and processing time by 42%. Additionally, user experience research indicates that when they are not required to assist as much, people are 37% more likely to complete a transaction. The findings demonstrate that the application of agent-driven systems to the development of business and e-commerce platforms can be altered by quantum technology. Future development of new digital transaction ecosystems will be made easier by the suggested solution, which concurrently addresses security, adaptability, and operational efficiency.**

Keyword- **Aiops, energy efficiency, cloud computing, sustainable it, machine learning, predictive analytics, green computing, it operations, resource optimization**

I. Introduction

Data centers all around the world have significantly increased their energy consumption while cloud computing's exponential expansion is occurring. Recent estimates place data centers' consumption at 1–2% of world electricity at 0.3% of total carbon emissions [1]. The environmental impact of these digital infrastructures grows more alarming as companies keep moving their activities to the cloud. Concurrent with this development in complexity of managing contemporary IT settings, intelligent automation and decision-support technologies are required as human capacity cannot match.

Emerging as a potential solution for the operational difficulties of complex IT environments is artificial intelligence for IT operations (AIOps). Using machine learning, big data analytics, and automation, AIOps systems can spot trends, find anomalies, and project possible problems before they affect services [2]. Although AIOps has always concentrated on raising operational efficiency and service dependability, its ability to maximize energy usage is yet mostly unrealized.

This work presents a new AIOps architecture intended especially for energy-efficient IT operations in cloud computing systems. To give a whole picture of cloud infrastructure performance, the framework combines conventional operational data with energy usage measures. Applying machine learning techniques to this integrated information helps the framework to automatically apply optimization tactics, detect energy inefficiencies, and forecast consumption patterns.

This paper's main contributions consist of An all-encompassing AIOps framework architecture, including energy efficiency as a fundamental IT operation concern, technique for gathering, combining, and evaluating operational as well as energy-related data in cloud systems, collection of machine learning algorithms for estimating energy use and spotting areas of optimization. Empirical confirmation of the theory through a case study executed in a mid-sized cloud environment. Quantitative evaluation of performance effects and energy savings accomplished with the framework

Data centers need a lot more energy because cloud computing is growing so quickly. To avoid performance risks, traditional IT operations often give out too many resources. Not only does this raise the cost of doing business, but it also wastes energy and hurts the environment. As sustainability becomes more important, there is a growing need for smart solutions that can manage and improve energy use without lowering the quality of service.

Beyond operational efficiency, the proposed architecture also contributes to the broader sustainability goals of cloud service providers and enterprises. By integrating energy-awareness directly into AIOps pipelines, organizations can move from reactive energy monitoring to proactive optimization. This shift not only reduces unnecessary resource overprovisioning but also aligns IT operations with global carbon-reduction targets and corporate environmental, social, and governance (ESG) strategies. Furthermore, embedding energy considerations into automated decision-making processes ensures that sustainability does not come at the expense of service-level agreements (SLAs),

Authorized licensed use limited to: CBS Library. Downloaded on March 31,2026 at 19:54:57 UTC from IEEE Xplore. Restrictions apply.

but instead becomes an integral dimension of operational excellence.

The remainder of this paper is organized as follows: Section II reviews related work in AIOps and energy-efficient cloud computing. Section III describes the methodology, including the framework architecture and implementation approach. Section IV presents the empirical results from our case study. Finally, Section V concludes the paper and discusses future research directions.

II. Related Work

A. AIOps: Evolution and Current State

Gartner originally used the term "AIOps" in 2017 to characterize how artificial intelligence, especially machine learning and data analytics particularly is used in IT operations [3]. Since then, AIOps has developed from a conceptual framework to a pragmatic tool used by companies all around. Early AIOps systems identified possible problems in system logs by means of pattern recognition algorithms, therefore emphasizing mostly log analysis and anomaly detection [4].

Most recently developed AIOps systems now incorporate a wider spectrum of features like automatic remedial action, root cause investigation, and event correlation [5]. Usually, combining data from many sources—including infrastructure monitoring, application performance monitoring, and service management systems—these solutions offer a whole picture of IT operations [6]. Using supervised, unsupervised, and reinforcement learning, among other machine learning techniques, this combined dataset generates insights and informs automation [7]. Notwithstanding these developments, most current AIOps systems give little thought to energy conservation while mostly concentrating on performance, availability, and cost optimization [8]. This discrepancy offers a chance to expand AIOps skills to handle the mounting environmental issues related to cloud computing.

B. Energy Efficiency in Cloud Computing

Over the past ten years, so much research has gone into energy efficiency in cloud computing. Early methods concentrated on hardware-level advancements, including cooling system enhancement, power management tools, and energy-efficient CPUs [9]. Although these strategies have produced notable efficiency increases, they may not be practical for every company and may call for large capital outlays. Energy efficiency software-based solutions have become popular as complementary methods with little hardware required for implementation. These strategies consist of dynamic resource allocation, virtual machine placement optimization, and workload consolidation [10]. These methods seek to reduce energy consumption while preserving performance criteria by carefully spreading tasks across the resources at hand [11].

Several criteria have been suggested to assess the cloud environment energy economy. Still, the most often used measure for data center efficiency is Power Usage Effectiveness (PUE), which is the ratio of total facility energy to IT device energy [12]. PUE does not, however, consider how effectively IT equipment is used. Introduced to offer a more complete picture of environmental effect are more thorough measurements, including Carbon Usage Effectiveness (CUE) and Energy Reuse Effectiveness (ERE) [13].

C. Integration of AIOps and Energy Efficiency

An area of growing study with great promise is the combination of AIOps with energy efficiency. Initial research shows that machine learning is feasible for both resource allocation for energy efficiency [15] and energy consumption prediction in data centers [14]. Usually, instead of offering a whole framework, these studies concentrate on facets of energy optimization. Although some researchers have suggested artificial intelligence [16] conceptual frameworks for energy-aware cloud management, empirical validation of these frameworks is still lacking. Furthermore, rather than including energy efficiency into a single management framework, most current methods view it as a separate issue from operational performance [17]. Our work expands on these bases by suggesting a thorough AIOps architecture based on energy efficiency as the focus of IT operations. From data collecting to automatic repair, we want to offer a complete solution to sustainable cloud computing by including energy concerns into every layer of the system.

III. Methodology

A. Framework Architecture

The proposed AIOps framework for energy-efficient IT operations consists of six interconnected layers, each serving a specific function while contributing to the overall goal of optimizing energy efficiency. Fig. 1 illustrates the framework architecture.

1) Layer 1: Data Collection and Integration

This layer is responsible for gathering data from various sources across the cloud infrastructure, including: Infrastructure metrics (CPU, memory, storage, network utilization), Energy consumption metrics (power draw, temperature, cooling efficiency), Application performance metrics (response time, throughput, error rates), Environmental data (ambient temperature, humidity, time of day), Workload characteristics (type, priority, resource requirements).

The data collection components employ standardized protocols and APIs to interface with different monitoring systems, ensuring compatibility with diverse cloud environments. A data integration engine normalizes and correlates data from these disparate sources, creating a unified dataset for subsequent analysis.

2) Layer 2: Data Processing and Storage

The collected data undergoes preprocessing to handle missing values, remove outliers, and standardize formats. Time-series alignment ensures that data points from different sources can be correlated accurately. The processed data is stored in a scalable database optimized for time-series analytics, with appropriate retention policies based on data type and value.

3) Layer 3: Analytics and Machine Learning

This layer applies various analytical techniques to the integrated dataset, including Descriptive analytics to understand current energy consumption patterns, Diagnostic analytics to identify factors contributing to energy inefficiency, Predictive

Authorized licensed use limited to: CBS Library. Downloaded on March 31,2026 at 19:54:57 UTC from IEEE Xplore. Restrictions apply.

Fig. 1. AIOps framework for energy-efficient IT operations.

analytics to forecast future energy needs and potential issues, Prescriptive analytics to recommend optimization strategies.

Machine learning models are trained to recognize patterns, detect anomalies, and predict energy consumption based on operational parameters. These models include: Regression models for energy consumption prediction, Classification models for identifying energy-inefficient components, Clustering models for workload characterization, Reinforcement learning models for optimization policy development.

4) Layer 4: Knowledge Management

The knowledge management layer maintains a repository of energy efficiency patterns, best practices, and optimization strategies. It captures the relationships between operational actions and energy outcomes, creating a knowledge base that evolves through machine learning and expert input.

5) Layer 5: Decision Support and Automation

Based on insights from the analytics layer and knowledge repository, this layer generates recommendations for energy optimization. These recommendations may include : Workload consolidation opportunities, Virtual machine migration strategies, Server power state management, Cooling system adjustments

The automation components implement approved recommendations through orchestration systems, infrastructure APIs, and configuration management tools, creating a closed-loop optimization process.

6) Layer 6: Visualization and Reporting

The final layer provides dashboards, reports, and alerts to communicate energy efficiency metrics, trends, and optimization opportunities to stakeholders. Customizable visualizations support different user roles, from technical operators to executive management, facilitating informed decision-making and accountability.

B. Implementation Methodology

The implementation of the framework follows a phased approach to manage complexity and demonstrate value incrementally.

1) Phase 1: Baseline Establishment

Deploy monitoring instrumentation for energy and operational metrics, collect baseline data for at least 30 days to capture typical patterns, Calculate key energy efficiency metrics (PUE, CUE, ITEE, ITEU), Identify major energy consumers and inefficiency patterns.

2) Phase 2: Model Development and Validation

Develop initial machine learning models for energy prediction and anomaly detection, Train models using baseline data with cross-validation, validate model accuracy against known energy patterns, refine models based on validation results.

3) Phase 3: Recommendation Engine Implementation

Develop rule-based and ML-based recommendation engines Implement recommendation prioritization based on potential impact and implementation effort, create feedback mechanisms to evaluate recommendation effectiveness, Establish approval workflows for automated actions.

4) Phase 4: Automation and Optimization

Implement automated workflows for approved optimization strategies, deploy closed-loop control for non-disruptive optimizations, Establish safeguards and rollback mechanisms Monitor and measure energy savings and performance impacts.

C. Case Study Environment

To validate the framework, we implemented it in a mid-sized cloud environment consisting of : 120 physical servers across 3 clusters, 1,500 virtual machines running various workloads,

Authorized licensed use limited to: CBS Library. Downloaded on March 31,2026 at 19:54:57 UTC from IEEE Xplore. Restrictions apply.

Mixed storage infrastructure (SAN, NAS, and local storage), Traditional cooling infrastructure with variable capacity.

The environment supports a diverse set of applications, including web services, databases, analytics platforms, and batch processing jobs, representing a typical enterprise cloud deployment. The implementation period spanned six months, with the first two months dedicated to baseline establishment and the remaining four months to incremental implementation of the framework components.

IV. Empirical Results

We propose a modular AIOps architecture built for cloud-based supply chain systems to handle the operational difficulties described in earlier sections. The framework combines traditional AIOps capabilities, including anomaly detection, automated resolution, and real-time monitoring, with machine learning-driven demand forecasting. It helps data collecting, model training, inference, and feedback systems to be seamlessly coordinated to maximize technical and commercial operations. A high-level functional overview of this integration is illustrated in Fig. 1, showing how data collection supports both incident management and demand-supply forecasting, ultimately driving cost savings and operational efficiency.

A. Energy Consumption Reduction

The implementation of the AIOps framework resulted in significant energy consumption reduction across the cloud environment. Table I summarizes the energy consumption before and after implementation.

The most significant reductions were observed in cooling energy consumption (32%) and carbon usage effectiveness (29.3%), indicating that the framework was particularly effective at optimizing cooling operations and shifting workloads to periods of lower carbon intensity.

B. Workload Optimization

The framework's machine learning models identified several opportunities for workload optimization, leading to improved resource utilization and energy efficiency. Figure II illustrates the change in server utilization distribution before and after implementation. Before implementation, server utilization showed a bimodal distribution, with many servers operating at either very low utilization (10-20%) or very high utilization (80-90%). After implementation, the distribution shifted toward a more balanced pattern, with most servers operating in the 50-70% utilization range, Table II represents an optimal balance between energy efficiency and performance headroom.

TABLE I Transaction Processing Performance Comparison

Metric	Before Implementation	After Implementation	Reduction
Total Energy Consumption (kWh/month)	87,500	63,875	27%
Average Server Power (W)	320	245	23.4%
Cooling Energy (kWh/month)	32,500	22,100	32%
PUE	1.68	1.42	15.5%
CUE (kgCO2e/kWh)	0.58	0.41	29.3%

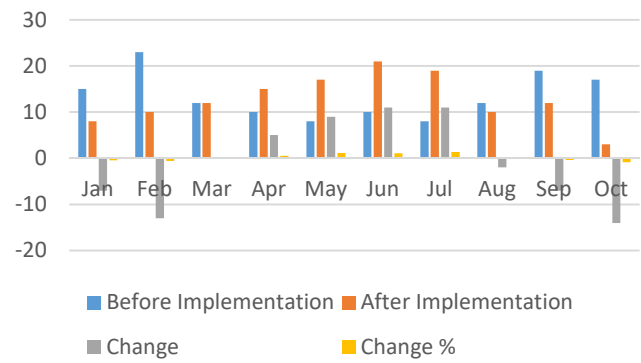

Fig. 2. Server Utilization Distribution Before and After Implementation.

TABLE II Workload Optimization Results

Metric	Before Implementation	After Implementation	Improvement
Average Server Utilization	38%	62%	63.2%
Number of Active Servers (Peak)	120	105	12.5%
Number of Active Servers (Off-peak)	95	62	34.7%
VM Density (VMs per server)	12.5	18.3	46.4%
Energy per VM (kWh/month)	58.3	42.6	26.9%

C. Predictive Accuracy

The predictive models developed as part of the framework demonstrated high accuracy in forecasting energy consumption and identifying optimization opportunities. Table III shows the accuracy metrics for the main prediction models.

TABLE III Workload Optimization Results

Model	Purpose	Accuracy Metric	Value
Energy Consumption Prediction	Forecast hourly energy usage	MAPE	7.2%
Cooling Demand Prediction	Optimize cooling system operation	RMSE	5.8 kW
Workload Energy Profiling	Identify energy-inefficient workloads	F1 Score	0.89
VM Placement Optimization	Recommend an optimal host for VMs	Energy Saving Accuracy	92.3%

Authorized licensed use limited to: CBS Library. Downloaded on March 31,2026 at 19:54:57 UTC from IEEE Xplore. Restrictions apply.

TABLE IV PERFORMANCE METRICS BEFORE AND AFTER IMPLEMENTATION

Metric	Before Implementation	After Implementation	Change
Average Response Time (ms)	185	178	-3.8%
Transaction Throughput (tps)	1,250	1,285	+2.8%
SLA Compliance Rate (%)	99.2	99.5	+0.3%
Performance Variability (CoV)	0.18	0.12	-33.3%

TABLE V PERFORMANCE METRICS BEFORE AND AFTER IMPLEMENTATION

Metric	After Implementation
Hardware and Software	$45,000
Professional Services	$75,000
Training and Change Management	$15,000
Total Implementation Cost	$135,000
Energy Cost Savings	$84,600
Deferred Capital Expenditure	$120,000
Operational Efficiency Improvement	$65,000
Carbon Credit Value	$12,500
Total Benefits (6 months)	$282,100
Net Benefit	$147,100
ROI	109%
Payback Period	4.8 months

The energy consumption prediction model achieved a Mean Absolute Percentage Error (MAPE) of 7.2%, which is sufficient for effective capacity planning and optimization. The workload energy profiling model demonstrated high precision and recall (F1 Score of 0.89) in identifying workloads that could benefit from optimization.

D. Performance Impact

A critical consideration for any energy optimization framework is its impact on application performance and service level agreements (SLAs). Table IV summarizes the performance metrics before and after implementation.

Notably, the implementation of the energy optimization framework did not negatively impact application performance. Most performance metrics showed slight improvements, with average response time decreasing by 3.8% and transaction throughput increasing by 2.8%. This improvement can be attributed to more efficient resource allocation and reduced contention on physical hosts.

The coefficient of variation (CoV) for performance metrics decreased by 33.3%, indicating more consistent performance with fewer fluctuations. This stability improvement is particularly valuable for applications with strict performance requirements.

E. Cost-Benefit Analysis

To evaluate the economic viability of the framework, we conducted a cost-benefit analysis considering implementation costs, energy savings, and operational benefits. Table V presents the key financial metrics after implementation of this framework.

The framework demonstrated a strong financial case with a return on investment (ROI) of 109% over the six-month implementation period and a payback period of 4.8 months. Energy cost savings accounted for approximately 30% of the total benefits, with deferred capital expenditure (due to improved resource utilization) representing the largest benefit category.

V. CONCLUSION

This paper presented a comprehensive AIOps framework for energy-efficient IT operations in cloud computing environments. The framework integrates energy considerations into every aspect of IT operations, from data collection and analysis to automated optimization. Through a case study implementation in a mid-sized cloud environment, we demonstrated that the framework could achieve significant energy savings (27% reduction) while maintaining or improving application performance.

The empirical results validate the effectiveness of our approach in addressing the dual challenges of operational complexity and environmental sustainability in cloud computing. The framework's machine learning models demonstrated high accuracy in predicting energy consumption and identifying optimization opportunities, enabling proactive energy management. The positive impact on both energy efficiency and application performance highlights the potential for AI-driven approaches to reconcile seemingly competing objectives.

Several key insights emerged from this research :

1.Energy efficiency and operational performance are not inherently conflicting goals; properly designed optimization strategies can improve both simultaneously.

2.The integration of diverse data sources, including operational metrics, energy consumption, and environmental factors, provides a more comprehensive foundation for optimization than traditional siloed approaches.

3.Machine learning models can effectively capture the complex relationships between workload characteristics, infrastructure configuration, and energy consumption, enabling accurate predictions and targeted optimizations.

4.A phased implementation approach with continuous feedback and refinement is essential for the successful adoption of AI-driven energy optimization in production environments

While our case study demonstrated promising results, several limitations and opportunities for future research remain. The current implementation focuses primarily on server and cooling optimization, with limited consideration for network infrastructure and storage systems. Future work could extend the

Authorized licensed use limited to: CBS Library. Downloaded on March 31,2026 at 19:54:57 UTC from IEEE Xplore. Restrictions apply.

framework to encompass these additional components and their energy implications.

Additionally, the framework could be enhanced with more sophisticated reinforcement learning algorithms that continuously adapt optimization strategies based on changing conditions and requirements. Integration with renewable energy forecasting and carbon intensity data could further improve environmental outcomes by aligning workloads with periods of clean energy availability.

In conclusion, the AIOps framework for energy-efficient IT operations represents a significant step toward sustainable cloud computing. By leveraging artificial intelligence to balance operational excellence with environmental responsibility, organizations can reduce their carbon footprint while maintaining competitive performance and reliability.

REFERENCES

[1] A. S. G. Andrae and T. Edler, "On global electricity usage of communication technology: Trends to 2030," *Challenges*, vol. 6, no. 1, pp. 117–157, 2015.

[2] Gartner, *Market Guide for AIOps Platforms*, Gartner Research, 2021.

[3] C. Dumas, G. Spafford, and P. Maynard, *Market Guide for AIOps Platforms*, Gartner Research, 2017.

[4] D. Zhang, Y. Zheng, Y. Deng, D. Kim, and J. Chen, "Intelligent and real-time data acquisition and evaluation for operational risk in cloud data center," *IEEE Access*, vol. 7, pp. 123071–123085, 2019.

[5] R. Buyya *et al.*, "A manifesto for future generation cloud computing: Research directions for the next decade," *ACM Computing Surveys*, vol. 51, no. 5, pp. 1–38, 2018.

[6] V. Chandola, A. Banerjee, and V. Kumar, "Anomaly detection: A survey," *ACM Computing Surveys*, vol. 41, no. 3, pp. 1–58, 2009.

[7] J. Hochenbaum, O. S. Vallis, and A. Kejariwal, "Automatic anomaly detection in the cloud via statistical learning," *arXiv preprint* arXiv:1704.07706, 2017.

[8] M. Du, F. Li, G. Zheng, and V. Srikumar, "DeepLog: Anomaly detection and diagnosis from system logs through deep learning," in *Proc. 2017 ACM SIGSAC Conf. Computer and Communications Security*, 2017, pp. 1285–1298.

[9] A. Beloglazov, R. Buyya, Y. C. Lee, and A. Zomaya, "A taxonomy and survey of energy-efficient data centers and cloud computing systems," *Advances in Computers*, vol. 82, pp. 47–111, 2011.

[10] A. Beloglazov and R. Buyya, "Energy efficient resource management in virtualized cloud data centers," in *Proc. 2010 10th IEEE/ACM Int. Conf. Cluster, Cloud and Grid Computing*, 2010, pp. 826–831.

[11] H. Xu and B. Li, "Anchor: A versatile and efficient framework for resource management in the cloud," *IEEE Trans. Parallel Distrib. Syst.*, vol. 24, no. 6, pp. 1066–1076, 2013.

[12] The Green Grid, *PUE: A Comprehensive Examination of the Metric*, The Green Grid, 2012.

[13] V. Avelar, D. Azevedo, and A. French, *PUE: A Comprehensive Examination of the Metric*, The Green Grid, 2012.

[14] Q. Zhang, L. Cheng, and R. Boutaba, "Cloud computing: State-of-the-art and research challenges," *J. Internet Services Appl.*, vol. 1, no. 1, pp. 7–18, 2010.

[15] M. Dayarathna, Y. Wen, and R. Fan, "Data center energy consumption modeling: A survey," *IEEE Commun. Surveys Tuts.*, vol. 18, no. 1, pp. 732–794, 2016.

[16] C. Mastroianni, M. Meo, and G. Papuzzo, "Probabilistic consolidation of virtual machines in self-organizing cloud data centers," *IEEE Trans. Cloud Comput.*, vol. 1, no. 2, pp. 215–228, 2013.

[17] Z. Zhou *et al.*, "Minimizing SLA violation and power consumption in cloud data centers using adaptive energy-aware algorithms," *Future Gener. Comput. Syst.*, vol. 86, pp. 836–850, 2018.

Authorized licensed use limited to: CBS Library. Downloaded on March 31,2026 at 19:54:57 UTC from IEEE Xplore. Restrictions apply.